

---

# Lauren Zhu

laurenzhu@gmail.com | [linkedin.com/in/laurenzhu](https://www.linkedin.com/in/laurenzhu) | [heylaurenz.github.io](https://heylaurenz.github.io) | (858) 380-9511

## BIO

My expertise is in AI and product quality—building products that ✨*work like magic*✨. I'm super plugged into the AI community and latest tech/trends. I'm endlessly curious, product minded, and want to build products that win and stay winners. Talk to me about context engineering, LLM as a judge, or the future of Agents! 🤖👨‍🔬

## EDUCATION

**Stanford University, CA — M.S.** in Computer Science, AI Concentration 2020 - 2021 (3.87)

**Stanford University, CA — B.S.** in Computer Science, AI Concentration 2016 - 2020 (3.8)

- 4x TA/Head TA for popular AI and NLP courses; President of Women in CS

## SKILLS

Python, Go, LLM stuff (prompt/context engineering, infra), tool orchestration, MCP, RAG, OpenAI Agents SDK, Claude Code, GCP products, Beam, obviously Cursor

## EXPERIENCE

### Glean — Senior Software Engineer — Aug 2021 - Present

- Tech Lead on [Agents Quality](#) (5 ppl): Built AI stack for the initial product; lead the team to scale it. Now we incorporate AI in new ways (following technology trends) to increase growth + hill climb on LLM judge metrics.
- Doc Q&A: Area owner for (built + scaling) the primary tools that fetch/manipulate/return doc context (RAG sidekick) in all our LLM-powered use cases.
- Designed and built unified context selection (RAG, Agentic search via filesystem) to be used by all Agentic products.
- Core [Assistant Quality](#) contributor: Built many new features + guided product / quality decisions all around the Agentic and pre-Agentic engines of our famous Glean Assistant.
  - Invented critical LLM judge metric used to measure + hill climb on hard queries
- Query Understanding: Developed early query parsing systems, spellcheck (15,000 triggers/day), acronym glossary (10,000 triggers/day)
- [Patented] [Expert detection](#): System to determine who knows the most about any given topic/area. Core component in Glean's Knowledge Graph.
- Consistently in the SWE top 5th percentile on PR output for the last few years 💪⚡

### Apple Special Projects — ML Intern — Jun 2020 - Sep 2020

- Deep Learning — diagnosed/implemented data sampling processes for model training like sharding and upsampling on low-resourced classes

### Ford Greenfield Labs — ML Intern — Jun 2019 - Sep 2019

- Deep learning for autonomous vehicles, hackathon winner

### The University of Edinburgh, Scotland — AI Researcher — Jun 2018 - Aug 2018

- Zero shot neural machine translation in Tensorflow
- Used GANs to universalize language representations during training